

DA_HW5

Howard

2025-04-02

Problem 1

Prepare dataset

- Convert the 400 images into a 400×2576 data matrix.

```
library(png)
image_dir <- "D:/DA_HW/DAH/URL Faces"
image_files <- list.files(image_dir, pattern = "\\*.png$", full.names = TRUE)
n <- length(image_files)
if(n != 400){
  stop("Expected 400 images, but found ", n)
}
data_matrix <- matrix(NA, nrow = n, ncol = 46 * 56)
for (i in 1:n) {
  img <- readPNG(image_files[i])
  data_matrix[i, ] <- as.vector(t(img))
}
dim(data_matrix)
```

```
## [1] 400 2576
```

- Add an additional column indicating the physical gender label (same in HW2) .

```
subject_ids <- rep(1:40, each = 10)
subject_gender <- c(0,rep(1,6),0,1,0,1,0,rep(1,19),0,rep(1,8))
gender_labels <- subject_gender[subject_ids]
final_data <- cbind(data_matrix, gender = gender_labels)
dim(final_data)
```

```
## [1] 400 2577
```

(a) Identify the value of λ that minimum MSE in both LASSO and Ridge regression models.

- Using cross-validation to get λ

```
library(glmnet)
```

```
## 載入需要的套件：Matrix
```

```
## Loaded glmnet 4.1-8
```

```
# Split to X and y
X <- as.matrix(final_data[,-2577])

# y as an binary
y_bin <- as.matrix(final_data[,2577])

# Set up cross-validation for LASSO
cv_lasso <- cv.glmnet(X, y_bin, alpha = 1, family = "binomial", type.measure = "mse")
lambda_min_lasso <- cv_lasso$lambda.min

# Set up cross-validation for Ridge
cv_ridge <- cv.glmnet(X, y_bin, alpha = 0, family = "binomial", type.measure = "mse")
lambda_min_ridge <- cv_ridge$lambda.min

# Print optimal Lambdas
cat("Lasso optimal lambda:", lambda_min_lasso, "\n")
```

```
## Lasso optimal lambda: 0.006514476
```

```
cat("Ridge optimal lambda:", lambda_min_ridge, "\n")
```

```
## Ridge optimal lambda: 1.061697
```

Compare the selected features

```
# selected features
lasso_coef <- coef(cv_lasso, s = "lambda.min")
ridge_coef <- coef(cv_ridge, s = "lambda.min")
selected_features_lasso <- which(lasso_coef != 0)[-1] #exclude intercept
selected_features_ridge <- which(ridge_coef != 0)[-1]
cat("Numbers of selected features from lasso:",length(selected_features_lasso),"\n\n")
```

```
## Numbers of selected features from lasso: 70
```

```
cat("Numbers of selected features from ridge:",length(selected_features_ridge),"\n\n")
```

```
## Numbers of selected features from ridge: 2576
```

Observation

- Ridge regression will keep all the features nonzero.
- Lasso regression will selected the important features from the full model.
- According to the HW2, stepwise selected 14 features while lasso selected 96 features.

The model performance

```
mse_lasso <- min(cv_lasso$cvm)
mse_ridge <- min(cv_ridge$cvm)
cat("lasso MSE:", mse_lasso,"\n")
```

```
## lasso MSE: 0.1006881
```

```
cat("ridge MSE:", mse_ridge,"\n")
```

```
## ridge MSE: 0.118241
```

Observation

- The MSE of ridge and lasso is almost the same.
- Ridge MSE has a little bit higher than lasso.

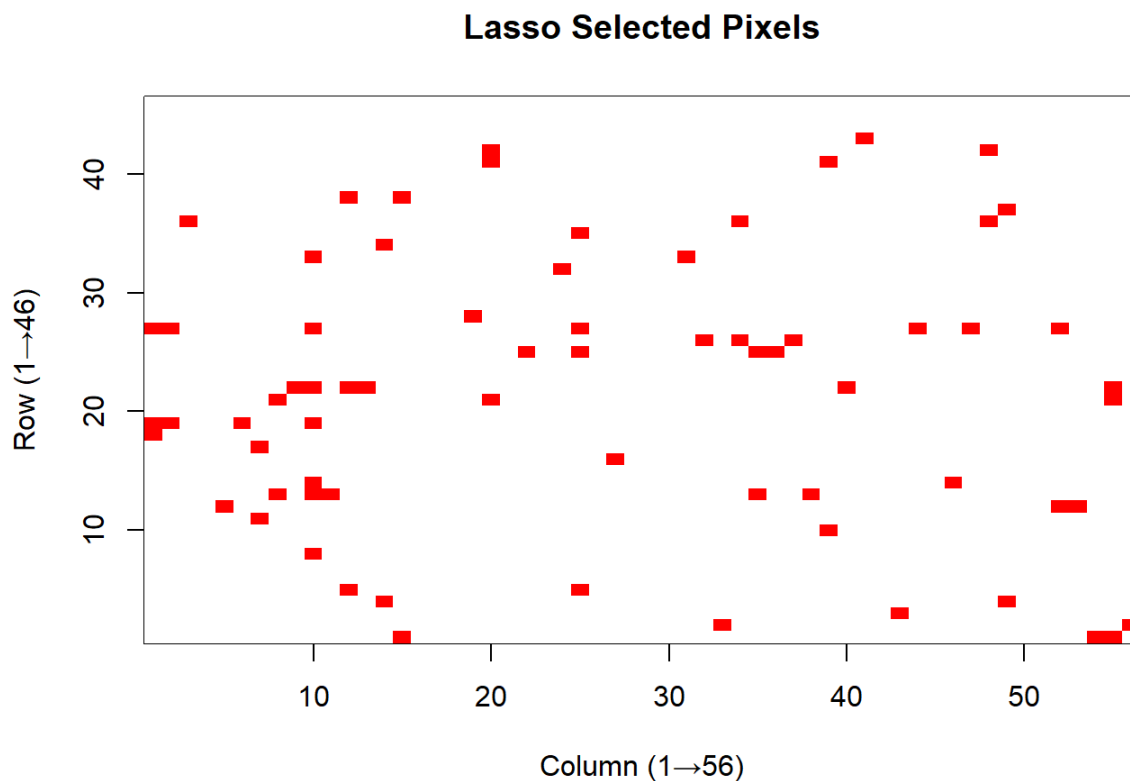
(b) Plot the chosen pixels from Lasso on 46×56 canvas.

```
# Initialize all zeros
important_pixel_map <- rep(0, 46 * 56)

# Mark selected pixel indices as 1
important_pixel_map[selected_features_lasso] <- 1

# Convert to a 46x56 matrix. We assume row-major flattening.
pixel_matrix <- matrix(important_pixel_map, nrow = 46, ncol = 56, byrow = TRUE)

# We'll use base R's image() function with two colors:
image(1:56, 1:46,
      t(apply(pixel_matrix, 2, rev))), # flip to put row=1 at top
      col = c("white", "red"),
      xlab = "Column (1→56)",
      ylab = "Row (1→46)",
      main = "Lasso Selected Pixels")
```



Problem 2(a): Estimate β_1 and β_2 (unconstrained Cobb-Douglas)

- The model is:

$$V_t = \alpha K_t^{\beta_1} L_t^{\beta_2} \eta_t$$

- Taking logs:

$$\log(V_t) = \log(\alpha) + \beta_1 \log(K_t) + \beta_2 \log(L_t) + \log(\eta_t)$$

- So we can fit this using linear regression.

```
# Data
year <- 72:86
capital <- c(1209188, 1330372, 1157371, 1070860, 1233475, 1355769, 1351667, 1326248, 1089545, 1111942, 988165, 1069651, 1191677, 1246536, 1281262)
labor <- c(1259142, 1371795, 1263084, 1118226, 1274345, 1369877, 1451595, 1328683, 1077207, 1056231, 947502, 1057159, 1169442, 1195255, 1171664)
value_added <- c(11150.0, 12853.6, 10450.8, 9318.3, 12097.7, 12844.8, 13309.9, 13402.3, 8571.0, 8739.7, 8140.0, 10958.4, 10838.9, 10030.5, 10836.5)

# Log-transform
log_K <- log(capital)
log_L <- log(labor)
log_V <- log(value_added)

# Fit unconstrained linear model
model_unconstrained <- lm(log_V ~ log_K + log_L)
summary(model_unconstrained)
```

```
##
## Call:
## lm(formula = log_V ~ log_K + log_L)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.090518 -0.043559 -0.009005  0.022608  0.180405
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -9.6259      2.8996  -3.320  0.00611 **
## log_K         0.5057      0.5061   0.999  0.33743
## log_L         0.8455      0.4216   2.006  0.06799 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.07495 on 12 degrees of freedom
## Multiple R-squared:  0.8184, Adjusted R-squared:  0.7881
## F-statistic: 27.03 on 2 and 12 DF,  p-value: 3.591e-05
```

Results:

- $\hat{\beta}_1$ is 0.5057
- $\hat{\beta}_2$ is 0.8455

Problem 2(b): Re-estimate with constraint $\beta_1 + \beta_2 = 1$

- Substitute

$$\begin{aligned}\log(V_t) &= \log(\alpha) + \beta_1 \log(K_t) + (1 - \beta_1) \log(L_t) \\ \Rightarrow \log(V_t) &= \log(\alpha) + \beta_1 (\log(K_t) - \log(L_t)) + \log(L_t) \\ Z_t &= \log(K_t) - \log(L_t)\end{aligned}$$

```
# Create transformed variable Z = Log(K) - Log(L)
Z <- log_K - log_L

# Fit constrained model: Log(V) = a + b * Z + Log(L)
model_constrained <- lm(log_V - log_L ~ Z)
summary(model_constrained)
```

```
##
## Call:
## lm(formula = log_V ~ log_L ~ Z)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.12097 -0.07111  0.02166  0.04319  0.14354
##
## Coefficients:
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept) -4.712885    0.020753  -227.090  <2e-16 ***
## Z            0.009609    0.441507    0.022    0.983
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08017 on 13 degrees of freedom
## Multiple R-squared:  3.643e-05, Adjusted R-squared:  -0.07688
## F-statistic: 0.0004737 on 1 and 13 DF,  p-value: 0.983
```

```
# Recover  $\theta_1$  and  $\theta_2$ 
beta1_constrained <- coef(model_constrained)[2]
beta2_constrained <- 1 - beta1_constrained
alpha_constrained <- exp(coef(model_constrained)[1])

cat("Under beta_1 + beta_2 = 1 constraint:\n")
```

```
## Under beta_1 + beta_2 = 1 constraint:
```

```
cat("  alpha =", alpha_constrained, "\n")
```

```
##    alpha = 0.008978832
```

```
cat("  beta_1 =", beta1_constrained, "\n")
```

```
##    beta_1 = 0.009608932
```

```
cat("  beta_2 =", beta2_constrained, "\n")
```

```
##    beta_2 = 0.9903911
```

Problem3 (a): Create a PCA function in R as following:

Input:

- Data matrix x
- Boolean `isCorrMX`: if `TRUE`, use correlation matrix, if `FALSE`, use covariance matrix.

Output:

- Loading matrix (eigenvectors)
- Eigenvalue vector
- Score matrix(PCs)
- Screen plot

```

myPCA <- function(X, isCorrMX = FALSE) {
  # Center the data
  X_centered <- scale(X, center = TRUE, scale = FALSE)

  # Use correlation or covariance matrix
  S <- if (isCorrMX) cor(X_centered) else cov(X_centered)

  # Spectral decomposition
  eig <- eigen(S)
  eig_values <- eig$values
  eig_vectors <- eig$vectors

  # Scores (projected data)
  scores <- X_centered %*% eig_vectors

  # Scree plot
  var_explained <- eig_values / sum(eig_values)
  cum_var <- cumsum(var_explained)

  barplot(var_explained, main = "Scree Plot", xlab = "PC", ylab = "Proportion of Variance", ylim = c(0,1.005), col =
"lightblue")
  lines(x = 1:length(cum_var), y = cum_var, type = "b", col = "red", pch = 16)
  return(list(
    loadings = eig_vectors,
    eigenvalues = eig_values,
    scores = scores,
    explained_variance = var_explained,
    cumulative_variance = cum_var
  ))
}

```

Problem3 (b): Demonstrate using the AutoMPG dataset

- Choose continous features in `mtcars`


```

# Load AutoMPG
data <- mtcars[, c("mpg", "cyl", "disp", "hp", "drat", "wt", "qsec")]

# Apply PCA with covariance matrix
pca_cov <- myPCA(data, isCorrMX = FALSE)

```

Scree Plot



pca_cov\$loadings

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,]  0.038118683 -0.009182403  0.99018630 -0.06777104  0.091032234
## [2,] -0.012035357  0.003374958 -0.06293054  0.27013664  0.915542928
## [3,] -0.899595869 -0.435453668  0.03180796  0.00149054 -0.008470542
## [4,] -0.434786012  0.899788639  0.02307582 -0.02831634  0.001506989
## [5,]  0.002660185  0.003899832  0.04043028  0.03767426 -0.309878145
## [6,] -0.006239546 -0.004865696 -0.08407174 -0.17014740  0.028601267
## [7,]  0.006671139 -0.025018357 -0.07296532 -0.94406555  0.237860800
##           [,6]      [,7]
## [1,] -0.047546856  0.0534159080
## [2,] -0.195351163 -0.2157052870
## [3,]  0.003456672 -0.0031159748
## [4,]  0.002142719  0.0005991073
## [5,] -0.615423407 -0.7226017728
## [6,] -0.759760099  0.6211809371
## [7,]  0.059827262 -0.2063705962
```

pca_cov\$eigenvalues


```
## [1] 1.864032e+04 1.453913e+03 9.287215e+00 1.547302e+00 3.767684e-01
## [6] 9.632947e-02 8.008942e-02
```

pca_cov\$scores

##	[,1]	[,2]	[,3]	[,4]	[,5]
## Mazda RX4	79.604508	-2.185045	-2.0200517	2.245285237	0.01344313
## Mazda RX4 Wag	79.606653	-2.200296	-2.0823506	1.673220944	0.15393850
## Datsun 710	133.883622	5.086336	-2.2918334	0.006313295	-0.72062110
## Hornet 4 Drive	-8.526653	-44.943825	1.1925904	-0.581196178	0.19968597
## Hornet Sportabout	-128.690876	-30.782573	3.2980561	0.702524796	0.42806449
## Valiant	23.210972	-35.064450	-3.3305143	-1.055270888	0.46297227
## Duster 360	-159.302142	32.272157	0.6341391	0.112712468	-0.16253840
## Merc 240D	112.612037	-39.713494	-0.3729508	-0.632937272	-0.54440602
## Merc 230	103.530932	-7.508062	-1.5823514	-4.190056002	0.02702953
## Merc 280	67.053960	6.169342	-3.4630474	0.134640852	0.25971920
## Merc 280C	67.004596	6.167186	-4.8930874	-0.336919020	0.27499055
## Merc 450SE	-55.208115	10.389804	-1.6261498	0.122361324	1.07264106
## Merc 450SL	-55.170353	10.378191	-0.7209908	-0.069595605	1.19241779
## Merc 450SLC	-55.248045	10.387223	-2.8337717	-0.313410010	1.09782449
## Cadillac Fleetwood	-242.811053	-52.517193	-0.8968358	-0.640083233	-0.88270619
## Lincoln Continental	-236.365729	-38.290433	-1.0478970	-0.817051039	-0.82076238
## Chrysler Imperial	-224.733256	-16.112726	2.9650090	-1.163289413	-0.40598391
## Fiat 128	172.354040	-6.557351	5.6155723	-0.706234032	0.49064268
## Honda Civic	181.063169	-17.799738	3.3695791	0.849686753	-0.19320393
## Toyota Corolla	179.688451	-4.169901	6.8410068	-1.129472348	0.73851790
## Toyota Corona	121.217849	3.392122	-3.2223024	-1.352828440	-0.55179458
## Dodge Challenger	-80.162716	-34.957009	-1.7949185	1.678002929	0.54231083
## AMC Javelin	-67.575373	-28.866726	-2.5457469	1.300673985	0.61258531
## Camaro Z28	-150.347472	36.647349	-0.6444273	0.545177108	-0.42455965
## Pontiac Firebird	-164.658164	-48.208305	5.0263994	0.628392017	0.17516975
## Fiat X1-9	171.887607	-6.625607	0.6390338	0.223051855	-0.11932287
## Porsche 914-2	123.800069	-2.047783	1.3997952	1.720055078	-1.17571127
## Lotus Europa	137.075805	28.676071	5.4741583	0.654345796	-0.29439930
## Ford Pantera L	-159.413292	53.312861	2.4438242	0.830788451	-0.56427329
## Ferrari Dino	64.779716	62.867154	-2.2383697	1.340699650	-0.01717668
## Maserati Bora	-145.337438	138.970783	1.6312425	-1.388065608	0.13936792
## Volvo 142E	115.176688	13.833937	-2.9228092	-0.391523451	-1.00386179
##	[,6]	[,7]			
## Mazda RX4	-0.145743789	-0.01625996			
## Mazda RX4 Wag	-0.305979348	0.02657364			
## Datsun 710	0.330528872	0.06922414			
## Hornet 4 Drive	0.404866456	0.04689261			
## Hornet Sportabout	0.275590211	-0.21902795			
## Valiant	0.494446856	0.19290457			
## Duster 360	0.428496287	-0.13120571			
## Merc 240D	-0.157560830	0.38471761			
## Merc 230	0.031171567	-0.45211342			
## Merc 280	-0.531262993	-0.01310717			
## Merc 280C	-0.428801037	-0.21171180			
## Merc 450SE	-0.302070804	0.29420739			
## Merc 450SL	-0.074579089	0.08980607			
## Merc 450SLC	0.011212208	-0.07385653			
## Cadillac Fleetwood	-0.060680572	0.09179817			
## Lincoln Continental	-0.265583704	0.22570360			
## Chrysler Imperial	-0.612485078	0.39397464			
## Fiat 128	-0.283979549	0.23892019			
## Honda Civic	-0.364740047	-0.64850149			
## Toyota Corolla	-0.166834373	-0.07450829			
## Toyota Corona	0.508642847	-0.12598074			
## Dodge Challenger	0.399252213	0.08839916			
## AMC Javelin	0.215413070	-0.30735639			
## Camaro Z28	-0.109404694	-0.27275658			
## Pontiac Firebird	0.127255262	-0.02098970			
## Fiat X1-9	0.126781305	-0.08141744			
## Porsche 914-2	-0.117848272	0.06387658			
## Lotus Europa	0.527491729	0.43677211			
## Ford Pantera L	-0.031064521	-0.71341860			
## Ferrari Dino	0.004414121	0.49360236			


```
## Maserati Bora      0.106839003  0.16138855
## Volvo 142E        -0.033783308  0.06345039
```

```
# Apply PCA with correlation matrix
pca_cor <- myPCA(data, isCorrMX = TRUE)
```



pca_cor\$loadings

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,]  0.4127573 -0.08296098  0.2416477  0.766798834 -0.2127946 -0.09002238
## [2,] -0.4247315 -0.07844163  0.1880252  0.193926827  0.2383825  0.78055217
## [3,] -0.4225036  0.08239922 -0.1180359  0.587679498  0.1488730 -0.16169873
## [4,] -0.3877611 -0.33696384 -0.2027400 -0.006884691 -0.8314576  0.04607937
## [5,]  0.3311703 -0.44858845 -0.7552915  0.117073263  0.2217502  0.23319304
## [6,] -0.3913153  0.32236122 -0.4405532  0.107346915  0.1666473 -0.36612190
## [7,]  0.2399275  0.74932087 -0.2943885  0.061382684 -0.3278152  0.40735764
##           [,7]
## [1,] -0.35069828
## [2,] -0.27276548
## [3,]  0.63803744
## [4,] -0.03873894
## [5,]  0.03702814
## [6,] -0.61280387
## [7,]  0.13083413
```

pca_cor\$eigenvalues

```
## [1] 5.08609988 1.15656554 0.34485150 0.15793358 0.12949405 0.07585706 0.04919838
```

pca_cor\$scores

##	[,1]	[,2]	[,3]	[,4]	[,5]
## Mazda RX4	44.562221	5.1049302	16.413005	-40.76210	20.160326
## Mazda RX4 Wag	44.596795	5.6067520	16.135807	-40.70035	20.019244
## Datsun 710	75.333452	8.0928702	25.593362	-70.11810	24.928028
## Hornet 4 Drive	3.532565	15.9394935	4.422079	17.28800	33.605196
## Hornet Sportabout	-67.376666	0.7364498	-20.511654	74.98510	-3.356860
## Valiant	18.037186	15.9859066	8.437661	-4.56472	33.266069
## Duster 360	-96.650189	-23.3551977	-35.521914	71.07781	-60.200807
## Merc 240D	71.603636	22.9886462	27.025315	-45.77461	55.777970
## Merc 230	61.427492	13.5723823	19.634849	-50.49535	26.895661
## Merc 280	35.694548	3.1340790	11.527312	-37.56217	10.403748
## Merc 280C	35.260645	3.6998169	11.012373	-38.59885	10.504971
## Merc 450SE	-34.871862	-7.1718573	-12.471515	23.78602	-19.597152
## Merc 450SL	-34.319347	-7.2062608	-12.163121	24.45191	-19.910890
## Merc 450SLC	-35.109732	-6.7161964	-12.810364	22.87156	-19.586815
## Cadillac Fleetwood	-130.306592	1.9463352	-42.733399	134.46171	-9.922470
## Lincoln Continental	-129.197455	-2.5072956	-43.426793	127.35776	-19.926553
## Chrysler Imperial	-124.777831	-9.9948401	-43.089248	118.79202	-36.121932
## Fiat 128	102.474290	14.4827285	36.471586	-79.72312	40.721660
## Honda Civic	108.627425	17.8372102	39.076154	-82.94497	52.743461
## Toyota Corolla	106.984577	14.2107605	37.862345	-83.02879	39.931749
## Toyota Corona	68.363012	9.0129746	22.677295	-63.94761	23.212154
## Dodge Challenger	-41.454763	6.0535938	-10.955442	47.97465	11.833878
## AMC Javelin	-35.397953	5.0447507	-9.759137	39.58002	9.744851
## Camaro Z28	-92.874526	-24.5646654	-34.968317	64.49768	-61.175477
## Pontiac Firebird	-84.244897	4.1753754	-25.246651	98.91280	2.533799
## Fiat X1-9	100.209417	14.4180106	35.488320	-83.52092	41.994268
## Porsche 914-2	72.037257	7.7654241	25.523786	-60.49077	28.465885
## Lotus Europa	76.044504	-1.8454572	25.817101	-72.21014	5.169514
## Ford Pantera L	-99.426476	-32.2096540	-38.141323	66.80113	-77.060970
## Ferrari Dino	24.776965	-18.4712476	5.119313	-51.09724	-35.563273
## Maserati Bora	-106.520269	-59.6787624	-46.519448	36.28437	-143.484790
## Volvo 142E	62.962568	3.9129441	20.080662	-63.58273	13.995559
##	[,6]	[,7]			
## Mazda RX4	9.2405905	-43.774205			
## Mazda RX4 Wag	9.3753497	-43.857203			
## Datsun 710	16.1164263	-75.916033			
## Hornet 4 Drive	-5.8370290	18.608089			
## Hornet Sportabout	-18.5828347	81.119324			
## Valiant	-0.2808766	-1.156083			
## Duster 360	-15.4754662	79.718844			
## Merc 240D	8.4965793	-50.941398			
## Merc 230	12.3648734	-55.010639			
## Merc 280	9.2267343	-39.058494			
## Merc 280C	9.5971802	-38.489016			
## Merc 450SE	-4.6248696	27.670171			
## Merc 450SL	-4.4999368	27.589063			
## Merc 450SLC	-4.1662528	28.347223			
## Cadillac Fleetwood	-34.8864452	153.336424			
## Lincoln Continental	-32.5978256	145.167616			
## Chrysler Imperial	-29.1401416	130.322375			
## Fiat 128	19.1937378	-96.736712			
## Honda Civic	19.2391731	-97.141411			
## Toyota Corolla	20.5829805	-101.787989			
## Toyota Corona	14.9434524	-67.806071			
## Dodge Challenger	-12.8367394	56.329369			
## AMC Javelin	-10.2487211	47.624842			
## Camaro Z28	-13.9212129	73.486706			
## Pontiac Firebird	-25.2481772	106.218620			
## Fiat X1-9	19.4691707	-94.668922			
## Porsche 914-2	13.1704024	-69.231041			
## Lotus Europa	18.0199807	-87.318957			
## Ford Pantera L	-13.4435886	72.821622			
## Ferrari Dino	14.2669843	-55.634779			

## Maserati Bora	-2.2792829	38.192626
## Volvo 142E	14.7657847	-68.023962

Is PCA scale-invariant ?

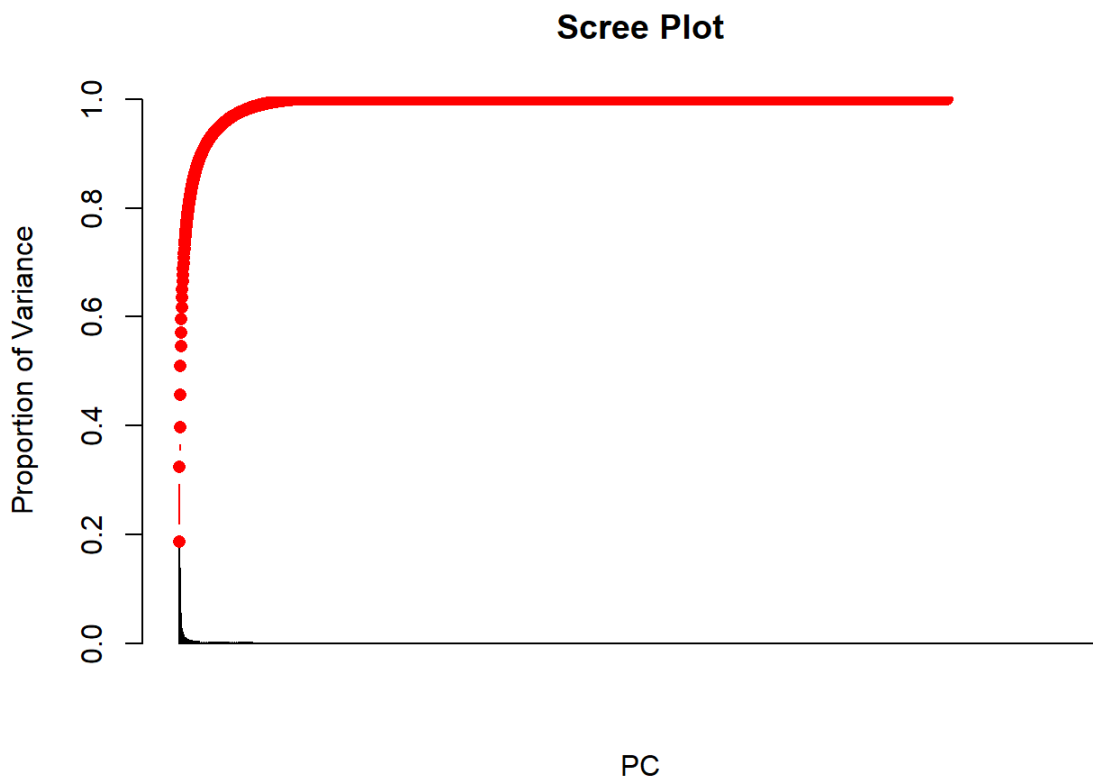
- PCA is not scale-invariant when using covariance matrix, because features with large variance dominate the PCs.
- Using the correlation matrix makes PCA scale-invariant, giving equal weight to each feature.

Problem 4 (a) :

- Run PCA, and determine the number of PCs

```
X <- t(data_matrix)

# Apply custom PCA
pca_result <- myPCA(t(X), isCorrMX = FALSE)
```



```
# Extract cumulative variance
cum_var <- pca_result$cumulative_variance

# Thresholds
thresholds <- c(0.5, 0.6, 0.7, 0.8, 0.9)

# Find number of PCs for each threshold
num_pcs_needed <- sapply(thresholds, function(th) which(cum_var >= th)[1])

# Print results
thresholds_percent <- thresholds * 100
for (i in seq_along(thresholds)) {
  cat(sprintf("To explain at least %d%% variance, need %d PCs\n", thresholds_percent[i], num_pcs_needed[i]))
}
```

```
## To explain at least 50% variance, need 5 PCs
## To explain at least 60% variance, need 9 PCs
## To explain at least 70% variance, need 16 PCs
## To explain at least 80% variance, need 32 PCs
## To explain at least 90% variance, need 76 PCs
```

Problem 4 (b):

- Rescale the first PC to the range [0,255], reshape it into a 46×56 matrix, and visualize it as a grayscale image using the scaled PC score.

```
# Extract first loading (eigenvector)
pc1 <- pca_result$loadings[, 1]

# Scale to [0, 255]
pc1_scaled <- 255 * (pc1 - min(pc1)) / (max(pc1) - min(pc1))

# Reshape to 46 x 56
pc1_matrix <- matrix(pc1_scaled, nrow = 46, ncol = 56)
```

- Use `ggplot` to visualize

```
# Visualize grayscale
# Step 5: High-quality grayscale visualization using ggplot2
library(ggplot2)
```

```
## Warning: 套件 'ggplot2' 是用 R 版本 4.2.3 來建造的
```

```
library(reshape2)
```

```
## Warning: 套件 'reshape2' 是用 R 版本 4.2.3 來建造的
```

```
# Convert matrix to data frame for ggplot
pc1_df <- melt(pc1_matrix)
colnames(pc1_df) <- c("y", "x", "value")

# Invert y-axis to display image properly (top to bottom)
pc1_df$y <- max(pc1_df$y) - pc1_df$y + 1

# Plot with ggplot2
ggplot(pc1_df, aes(x = x, y = y, fill = value)) +
  geom_tile() +
  scale_fill_gradient(low = "black", high = "white") +
  theme_void() +
  coord_fixed() +
  ggtitle("First Principal Component (PC1) - 46 x 56 Grayscale Image") +
  theme(plot.title = element_text(hjust = 0.5))
```

First Principal Component (PC1) - 46 × 56 Grayscale Image

